

## Week 4

This week, we will learn how to use agents with Jenkins and how to compile code.

Agents are the execution environment to run work in the pipeline. They may be a physical server or docker container running in a kubernetes pod. Last week, the agent was the same as the Jenkins master. Running on the master is unsafe- if the pipeline was to crash, so would the master. It also does not scale. If many people were to run the pipeline at once the master's CPU would be under pressure.

We will set up agents to be separate pods. That pod could be scheduled on any machine within the kubernetes cluster. If you were running kubernetes in a more realistic setting with multiple servers, rather than your laptop, the pod could be scheduled on any one of those servers. Jenkins could then scale up to support more users.

This week's configuration will also set up one form of persistent storage, so you will not lose your pipelines in certain failure cases.

The agent pods will run an image with the tools to compile Java. We will build a pipeline that uses those tools. The code will be given to you (you do not need to write Java code in this course!), but your pipeline must compile it whenever a change is made in GitHub. There will be a second stage to test the code.

The labs will be demoed in the chat session and recorded (as will be done every week!)

## Lab 1: setting up jenkins with multiple agents

We will tear down last week's Jenkins and build a new one to work with agents and kubernetes. Notice how simple this is using Kubernetes. This will clean up resources and is good practice.

First, delete everything in last week's work. Notice that by deleting the namespace, the deployment and its pods will also be deleted.

```
kubect1 delete all --all -n jenkins
```

Pull or clone week4's files from the umlF21 repository

```
git pull https://github.com/dlambrig/umlF21.git
```

We are interested in week4's yaml files describing the deployment and service.

Create a namespace, then the pvc, then then deployment. **Note we start everything in devops-tools namespace**

```
kubect1 create namespace devops-tools  
kubect1 apply -f jenkins-pvc.yaml -n devops-tools  
kubect1 apply -f jenkins-full-deployment.yaml -n devops-tools
```

This will start up Jenkins. Make sure Jenkins is running and you can get into it. Once that is done, follow instructions [here](#) with *customization* described in the next few pages. In the instructions, we will be running "Jenkins master and agents on the *same* kubernetes cluster".

Be careful,

- Confusingly, the instructions also show how to run "Jenkins master *outside* the cluster and agents in the cluster". Ignore those steps. In this lab's setup, both the jenkins master and agents run in Kubernetes.
- A mistake in setup can be hard to track down. If you feel a mistake was made, it may be easier to tear everything down and start from the beginning. Looking at jenkins logs is helpful for diagnosis. Setting pod retention to "always" in the kubernetes configuration will prevent agents from being deleted on failure, allowing diagnosis. The setup will be demoed and recorded.
- Apply the yaml files from the umlF21 git repository, not the one from the blog, in the devops-tools namespace. We won't use a certificate key or service account.

First, do steps 1 and 2 "Jenkins Kubernetes Plugin Configuration" per scenario 1 "Jenkins server running inside the same Kubernetes cluster"

## Configure Clouds

☒ **Kubernetes**

Name [?](#)  
kubernetes

Kubernetes URL [?](#)

☐ Use Jenkins Proxy [?](#)

Kubernetes server certificate key [?](#)

☒ Disable https certificate check [?](#)

Kubernetes Namespace  
devops-tools

Credentials  
- none - [Add](#)

☐ WebSocket [?](#)

☐ Direct Connection [?](#)

Jenkins URL [?](#)  
http://jenkins-service.devops-tools.svc.cluster.local:8080

[Test Connection](#)

Step 3 configures the "cloud" jenkins will connect with. In this case, its kubernetes. Make sure "test connection" works before continuing.

The namespace used in this demo is devops-tools.

Jenkins URL [?](#)  
http://jenkins-service.devops-tools.svc.cluster.local:8080

Jenkins tunnel [?](#)

Connection Timeout [?](#)  
5

Read Timeout [?](#)  
15

Concurrency Limit [?](#)  
10

Pod Labels [?](#)

**Pod Label**

Key [?](#)  
jenkins

Value [?](#)  
agent

[Delete Pod Label](#)

[Add Pod Label](#)

[Pod Retention...](#)

In step 4 configure the "jenkins URL": http://jenkins-service.devops-tools.svc.cluster.local:8080. This is how other pods will access jenkins. Our goal is to spawn agents (kubernetes pods) that will run our pipeline. Have a single "label" that will be associated with each pod. The label's value is "agent".

On Step 5, you create the "pod template". Important parts are "name", "namespace" and "labels".

#### Pod Templates

##### Pod Template

Name ?

kube-agent

Namespace ?

devops-tools

Labels ?

kubeagent

Usage ?

Use this node as much as possible

Pod template to inherit from ?

Containers ?

##### Container Template

Name ?

jnlp

Docker image ?

jenkins/inbound-agent:4.3-4

☐ Always pull image

Working directory ?

/home/jenkins/agent

Command to run ?

Add Container -

List of container in the agent pod

Environment variables ?

Add Environment Variable -

List of environment variables to set in all container of the pod

Volumes ?

Add Volume -

List of volumes to mount in agent pod

Annotations ?

Add Annotation -

List of annotations to set in agent pod

Concurrency Limit ?

Pod Retention ?

Never

Time in minutes to retain agent when idle ?

Time in seconds for Pod deadline ?

Timeout in seconds for Jenkins connection ?

1000

Raw YAML for the Pod ?

By the end of step 5, you have configured jenkins to run inside kubernetes and to spawn off pods to run pipelines.

In step 6, you create a simple "freestyle" pipeline which you should then be able to build. Check "console output" to verify the output is as expected, as described in the blog. For that simple pipeline, the console output should look like this:

```
Building remotely on kube-agent-rjjpk (kubeagent) in workspace  
/home/jenkins/agent/workspace/test  
[test] $ /bin/sh -xe /tmp/jenkins3564413985139030409.sh  
+ echo test  
test  
Finished: SUCCESS
```

Once this works, you are ready for a more complex pipeline in the next lab.

## Lab 2 Compiling Java Code

This section describes how to compile Java code in a github repository.

1. An agent will be spun up that contains the "maven" Java build tool.
2. Java code in a github repository shall be downloaded into the agent.
3. The maven build tool will compile the Java code.

Notice that because we use containers, all Java code is downloaded automatically. You will not need to install anything on your PC.

Here is the Jenkinsfile:

```
podTemplate(containers: [
  containerTemplate(
    name: 'maven',
    image: 'maven:3.8.1-jdk-8',
    command: 'sleep',
    args: '30d'
  ),
]) {

  node(POD_LABEL) {
    stage('Get a Maven project') {
      git 'https://github.com/dlambrig/simple-java-maven-
app.git'

      container('maven') {
        stage('Build a Maven project') {
          sh '''
          echo "maven build"
          mvn -B -DskipTests clean package
          '''
        }
      }
    }
  }
}
```

Now trying building the project, in the same manner as last week.

You are encouraged to download the blue ocean plugin, as described in the [Jenkins documentation](#). Once downloaded, click "blue ocean" on the main Jenkins page.

The blue ocean interface will allow you to observe the containers as they are being created. As you wait for the long process to complete, it is useful to know the machine is actually working (as opposed to being stuck).

The blue ocean interface will tell you when progress is made from the build stage to the test phase in a very nice display.